

Dependency management: Lessons learned

Simon Heybrock

`simon.heybrock@ess.eu`

European Spallation Source / DMSC / 2022-05-11

Context

- Developing Python and/or C++ libraries and applications
- Using GitHub (or similar) with Github Actions (or similar) to:
 - Build and run tests for pull requests (PRs).
 - Build documentation.
 - Build conda packages.
 - Build PyPI wheels.

Problem

Example: scipp, scippnexus, scippneutron, ess

- Experienced issues in about 10 cases in 2021, such as:
 - Pull request builds suddenly fail for no apparent reason.
 - Documentation builds stop working.
 - Package builds don't work any more.
- ...despite previously passing CI, issue is *not* a bad merge.

Previous approach

- Waste time discussing the “weird” issue at the daily standup.
- Dig through build logs: which package upgraded?
- Does older version work? Check upstream for issue, or open one.
- Pin version, restart all builds.
- Open issue as a reminder to remove the pin once upstream made a fix.
- Add pin across all projects.
- Remove pins once possible, half the time it is unclear how to reproduce.

Concrete examples from 2022

Five such issues is less than four months!

- February 18 Conan suddenly produces an error, caused by update of markupsafe <https://github.com/pallets/markupsafe/issues/282>. Lost ability to make a (patch) release until we found this and pinned the version.
- March 24 Documentation builds stop working due to <https://github.com/jupyter/nbconvert/issues/1736> causes by <https://github.com/pallets/jinja/issues/1626>. Lost ability to pass CI (blocking all merges) as well as blocking release builds (which build and update docs).
- April 1 We notice that navigation buttons in documentation pages have disappeared. The sphinx-book-theme 0.3.0 release causes this.
- April 7 scikit-build releases 0.14.0, breaking our PyPI builds. Lost ability to make a (patch) release until we found this and pinned the version.
- April 19 We notice that plots no longer show up in documentation pages, traced to an ipyml update <https://github.com/matplotlib/ipyml/issues/462>.

By continuously upgrading, we experience *almost every upstream issue*, even if it is fixed without hours or days.

Recommended reading

Approach

- Step 1: Read <https://hynek.me/articles/semver-will-not-save-you/>. Recommends to freeze everything, including transitive dependencies.
- Step 2: Conclude that we cannot do that because if we do that for a *library* almost no one will be able to use it.

Recommended reading

Approach

- Step 1: Read <https://hynek.me/articles/semver-will-not-save-you/>. Recommends to freeze everything, including transitive dependencies.
- Step 2: ~~Conclude that we cannot do that because if we do that for a *library* almost no one will be able to use it.~~
- Step 3: Realize that many or most of our dependency issues are not issues that ever reach a user. Actually we can freeze for quite a few cases:
 - Building documentation.
 - Requirements need for building packages.
 - Pull-request builds.
- Step 4: Must read for more details and better understanding:
<https://iscinumpy.dev/post/bound-version-constraints/>

Implementation

Requirement: All builds should be *reproducible*

Reproducible means:

- Builds for a commit that passed previously should continue passing if we re-run it later.
- Builds should fail *only* when there is a problem in the *current* source repository, *not* if a dependency or transitive dependency makes a release.

This requirement applies to “everything”:

- Going back to much older version, e.g., to make a patch release.
- Documentation builds (Sphinx).
- Release builds (making packages for PyPI and conda).
- PR builds.
- Builds on `main` after merging a PR.
- ...

Workflow

Setup (example for Python-based project with pip workflow)

- 1 Write `requirements.in` file(s). Do this for *everything* this is needed for your builds.
- 2 Use `pip-compile` from `pip-tools`. See also `pip-compile-multi`. Resulting `requirements.txt` is *checked into version control*.
- 3 Treat *warnings as errors* in CI, in particular `DeprecationWarning`
 - Avoids issues from *intentional* breaking changes, since we can take action early.

From time to time

- Re-run `pip-compile`, run all the builds (including release builds!) and check that things work, merge.
 - `dependabot` may be an alternative, but it updates individual requirements, which seems too noisy for us, with too little control.
- Make patch-releases when a dependency update causes issues for users
 - Much rarer than issues we experience in development.
 - Treating `DeprecationWarning` as errors should catch many issues months/years ahead.
 - Users can freeze dependency version *when setting up their environment* as a immediate workaround.

Status

- We (scipp and subprojects) do *not* have long term experience with this yet:
 - So far implemented in <https://github.com/scipp/scippnexus> and <https://github.com/scipp/scipp> (excluding C++ build dependencies).
 - Have not gone through update cycles (re-run of `pip-compile`) and do not know how many headaches this will cause.
- Our conan dependencies have been pinned since the beginning.
- conda dependencies are not pinned yet. What is best use of tools for keeping this up to date, including pins of transitive dependencies?
 - Does the conda-forge global pinning help? See https://conda-forge.org/docs/maintainer/pinning_deps.html.

Case study: Should we freeze dependencies in PR builds?

Temptation: *Use latest version of library dependencies to be alerted early of upstream changes*

Would this be a good idea?

Case study: Should we freeze dependencies in PR builds?

Temptation: *Use latest version of library dependencies to be alerted early of upstream changes*

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.

Case study: Should we freeze dependencies in PR builds?

Temptation: *Use latest version of library dependencies to be alerted early of upstream changes*

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.
- The builds do not pass, since a completely unrelated test is failing (because there was a unrelated dependency update). Alice is confused.

Case study: Should we freeze dependencies in PR builds?

Temptation: Use latest version of library dependencies to be alerted early of upstream changes

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.
- The builds do not pass, since a completely unrelated test is failing (because there was a unrelated dependency update). Alice is confused.
- Bob, the project maintainer, figures out that what the issue is, explains it in the PR, fixes the issue, merges into `main`, merges `main` into the PR, explains that all is good now, ...

Case study: Should we freeze dependencies in PR builds?

Temptation: Use latest version of library dependencies to be alerted early of upstream changes

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.
- The builds do not pass, since a completely unrelated test is failing (because there was a unrelated dependency update). Alice is confused.
- Bob, the project maintainer, figures out that what the issue is, explains it in the PR, fixes the issue, merges into `main`, merges `main` into the PR, explains that all is good now, ...
- Alice is frustrated by the delay and noise and moves on.

Case study: Should we freeze dependencies in PR builds?

Temptation: Use latest version of library dependencies to be alerted early of upstream changes

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.
- The builds do not pass, since a completely unrelated test is failing (because there was a unrelated dependency update). Alice is confused.
- Bob, the project maintainer, figures out that what the issue is, explains it in the PR, fixes the issue, merges into `main`, merges `main` into the PR, explains that all is good now, ...
- Alice is frustrated by the delay and noise and moves on.
- Bob is frustrated because he has to do things like this all the time.

Case study: Should we freeze dependencies in PR builds?

Temptation: Use latest version of library dependencies to be alerted early of upstream changes

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.
- The builds do not pass, since a completely unrelated test is failing (because there was a unrelated dependency update). Alice is confused.
- Bob, the project maintainer, figures out that what the issue is, explains it in the PR, fixes the issue, merges into `main`, merges `main` into the PR, explains that all is good now, ...
- Alice is frustrated by the delay and noise and moves on.
- Bob is frustrated because he has to do things like this all the time.

Case study: Should we freeze dependencies in PR builds?

Temptation: *Use latest version of library dependencies to be alerted early of upstream changes*

Would this be a good idea?

- A new contributor, Alice, opens a PR with a small bugfix.
- The builds do not pass, since a completely unrelated test is failing (because there was a unrelated dependency update). Alice is confused.
- Bob, the project maintainer, figures out that what the issue is, explains it in the PR, fixes the issue, merges into `main`, merges `main` into the PR, explains that all is good now, ...
- Alice is frustrated by the delay and noise and moves on.
- Bob is frustrated because he has to do things like this all the time.

Better solution: Freeze all dependencies for PR builds to *decouple* the tasks

- A new contributor, Alice, opens a PR with a small bugfix. Builds pass and it gets merged. Alice is happy and may continue contributing.
- *Independently*, Bob notices that a new upstream release requires a fix. He re-runs `pip-compile` and fixes the issue in a PR. After a merge, all future PR builds will now use the new upstream version.